

Assignment 3: Build Non-Linear Models – Notebook

This notebook contains:

- Applied Question 8 (a–f) from ISLR Python
- Two Abalone models (regularized regression and PCR) for Kaggle submission

Imports

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression, LassoCV, RidgeCV
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import KFold, cross_val_score
import statsmodels.api as sm
import itertools
import warnings
warnings.filterwarnings('ignore')

np.random.seed(123)
```

Applied Question 8(a): Generate X and noise ϵ

```
In [2]: # 8(a) Generate X and noise epsilon, n = 100
n = 100
X = np.random.normal(size=n)
eps = np.random.normal(scale=1.0, size=n)

print("First 10 values of X:")
print(X[:10])
print("\nFirst 10 values of eps:")
print(eps[:10])
```

First 10 values of X:
[-1.0856306 0.99734545 0.2829785 -1.50629471 -0.57860025 1.65143654
-2.42667924 -0.42891263 1.26593626 -0.8667404]

First 10 values of eps:
[0.64205469 -1.97788793 0.71226464 2.59830393 -0.02462598 0.03414213
 0.17954948 -1.86197571 0.42614664 -1.60540974]

Applied Question 8(b): Generate Y using cubic model

```
In [3]: # 8(b) Generate Y = β₀ + β₁ X + β₂ X² + β₃ X³ + ε
```

```
beta0 = 1
beta1 = 2
beta2 = 3
beta3 = 4

Y = beta0 + beta1 * X + beta2 * (X**2) + beta3 * (X**3) + eps

# create polynomial predictors X¹ ... X¹⁰
poly = PolynomialFeatures(degree=10, include_bias=False)
X_poly = poly.fit_transform(X.reshape(-1, 1))
poly_cols = [f"X{i}" for i in range(1, 11)]
df_poly = pd.DataFrame(X_poly, columns=poly_cols)

print("First 5 rows of polynomial predictor matrix:")
print(df_poly.head())
```

First 5 rows of polynomial predictor matrix:

	X1	X2	X3	X4	X5	X6	X7	\
0	-1.085631	1.178594	-1.279518	1.389083	-1.508031	1.637165	-1.777356	
1	0.997345	0.994698	0.992057	0.989424	0.986798	0.984178	0.981565	
2	0.282978	0.080077	0.022660	0.006412	0.001815	0.000513	0.000145	
3	-1.506295	2.268924	-3.417668	5.148015	-7.754428	11.680454	-17.594206	
4	-0.578600	0.334778	-0.193703	0.112076	-0.064847	0.037521	-0.021710	

	X8	X9	X10
0	1.929553	-2.094781	2.274159
1	0.978960	0.976361	0.973769
2	0.000041	0.000012	0.000003
3	26.502059	-39.919911	60.131151
4	0.012561	-0.007268	0.004205

Helper Functions for Stepwise Selection

```
In [4]: def calculate_rss(model, X, y):
    preds = model.predict(X)
    return np.sum((y - preds)**2)

def forward_stepwise(X, y):
    remaining = list(X.columns)
    selected = []
    models = []

    while remaining:
        rss_list = []
        for var in remaining:
            predictors = selected + [var]
            lm = LinearRegression().fit(X[predictors], y)
            rss = calculate_rss(lm, X[predictors], y)
            rss_list.append((rss, var, predictors, lm))
        best_rss, best_var, best_pred, best_model = min(rss_list, key=lambda
selected.append(best_var)
remaining.remove(best_var)
models.append((list(selected), best_rss, best_model))
```

```

    return models

def backward_stepwise(X, y):
    selected = list(X.columns)
    models = []

    while len(selected) > 0:
        rss_list = []
        for var in selected:
            trial = [v for v in selected if v != var]
            if not trial:
                lm = LinearRegression().fit(X[selected], y)
                rss = calculate_rss(lm, X[selected], y)
                rss_list.append((rss, None, list(selected), lm))
            else:
                lm = LinearRegression().fit(X[trial], y)
                rss = calculate_rss(lm, X[trial], y)
                rss_list.append((rss, var, list(trial), lm))
        best_rss, removed, best_pred, best_model = min(rss_list, key=lambda
models.append((list(best_pred), best_rss, best_model))
        if removed is None:
            break
        selected.remove(removed)
    return models

```

Applied Question 8(c): Forward Stepwise Selection

```
In [5]: # 8(c) Forward stepwise selection using X1 ... X10

fwd_models = forward_stepwise(df_poly, Y)

# choose model with smallest RSS as approximation to Cp-based choice
best_fwd = min(fwd_models, key=lambda x: x[1])
best_fwd_predictors, best_fwd_rss, best_fwd_model = best_fwd

print("Forward Stepwise Selected Predictors (approx Cp):")
print(best_fwd_predictors)
print("\nForward model coefficients (including intercept):")
print([best_fwd_model.intercept_] + list(best_fwd_model.coef_))

Forward Stepwise Selected Predictors (approx Cp):
['X3', 'X2', 'X1', 'X7', 'X10', 'X6', 'X9', 'X4', 'X8', 'X5']

Forward model coefficients (including intercept):
[numpy.float64(1.104588359466403), numpy.float64(4.046942427337782), numpy.float64(1.824954452936657), numpy.float64(2.1151709430889585), numpy.float64(0.031178989877454554), numpy.float64(-0.006519037015505686), numpy.float64(-0.6897976469084676), numpy.float64(-0.0019710196828176124), numpy.float64(1.5972609987753947), numpy.float64(0.1155541355619909), numpy.float64(-0.12376067397824232)]
```

Applied Question 8(d): Backward Stepwise Selection

```
In [6]: # 8(d) Backward stepwise selection using X1 ... X10

bwd_models = backward_stepwise(df_poly, Y)

# choose model with smallest RSS
best_bwd = min(bwd_models, key=lambda x: x[1])
best_bwd_predictors, best_bwd_rss, best_bwd_model = best_bwd

print("Backward Stepwise Selected Predictors (approx Cp):")
print(best_bwd_predictors)
print("\nBackward model coefficients (including intercept):")
print([best_bwd_model.intercept_] + list(best_bwd_model.coef_))
```

Backward Stepwise Selected Predictors (approx Cp):
['X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X7', 'X8', 'X10']

Backward model coefficients (including intercept):
[np.float64(1.0993139771027067), np.float64(2.174862012357655), np.float64(1.89226299313194), np.float64(3.8997233025200813), np.float64(1.4899470498957772), np.float64(-0.02384185483873066), np.float64(-0.6325842749254406), np.float64(0.00664902848955691), np.float64(0.10347195706955616), np.float64(-0.00564012540222601)]

Applied Question 8(e): Lasso with Cross Validation

```
In [7]: # 8(e) Lasso model using X1 ... X10

lasso = LassoCV(cv=10, random_state=123)
lasso.fit(df_poly, Y)


lasso_coefs = pd.Series(lasso.coef_, index=df_poly.columns)
print("Lasso selected coefficients:")
print(lasso_coefs)

plt.figure(figsize=(8,4))
plt.plot(lasso.alphas_, lasso.mse_path_.mean(axis=1))
plt.xlabel("Lambda")
plt.ylabel("Mean CV MSE")
plt.title("Lasso Cross-Validation Error Path")
plt.show()
```

Lasso selected coefficients:

X1	0.000000
X2	0.000000
X3	0.000000
X4	0.000000
X5	0.000000
X6	0.000000
X7	0.271679
X8	0.019014
X9	-0.020878
X10	0.000000

dtype: float64



Applied Question 8(f): Alternative Model $Y = \beta_0 + \beta_7 X^7 + \epsilon$

```
In [8]: # 8(f) New model: Y_alt = beta0 + beta7 * X**7 + eps

beta0_alt = 1
beta7_alt = 10

Y_alt = beta0_alt + beta7_alt * (X**7) + eps

# Forward stepwise on Y_alt
fwd_alt_models = forward_stepwise(df_poly, Y_alt)
best_fwd_alt = min(fwd_alt_models, key=lambda x: x[1])
best_fwd_alt_predictors, best_fwd_alt_rss, best_fwd_alt_model = best_fwd_alt

print("Forward Stepwise Selected Predictors for Y_alt:")
print(best_fwd_alt_predictors)

# Lasso on Y_alt
lasso_alt = LassoCV(cv=10, random_state=123)
lasso_alt.fit(df_poly, Y_alt)
lasso_alt_coefs = pd.Series(lasso_alt.coef_, index=df_poly.columns)

print("\nLasso Coefficients for Y_alt:")
print(lasso_alt_coefs)
```

```
Forward Stepwise Selected Predictors for Y_alt:  
['X7', 'X5', 'X1', 'X10', 'X8', 'X4', 'X6', 'X2', 'X9', 'X3']  
  
Lasso Coefficients for Y_alt:  
X1      0.000000  
X2      0.000000  
X3      0.000000  
X4      0.000000  
X5      0.000000  
X6      0.000000  
X7      0.000000  
X8      0.000000  
X9      1.287159  
X10     0.000000  
dtype: float64
```

Abalone Kaggle Models

This section builds two models using the Kaggle Abalone dataset:

- A regularized regression model (Ridge)
- A principal components regression (PCR) model

Both models:

- Use one-hot encoding for the Sex variable
- Use train.csv for training and test.csv for prediction
- Clip negative predictions to zero to satisfy MSLE
- Generate separate Kaggle submission files

Load train.csv and test.csv

```
In [9]: # Update these paths if your files are stored elsewhere  
train = pd.read_csv('/Users/abigailalbury-bloom/Downloads/train.csv')  
test = pd.read_csv('/Users/abigailalbury-bloom/Downloads/test.csv')  
  
print("Train shape:", train.shape)  
print("Test shape:", test.shape)  
print("Train columns:", train.columns.tolist())
```

```
Train shape: (90615, 10)  
Test shape: (60411, 9)  
Train columns: ['id', 'Sex', 'Length', 'Diameter', 'Height', 'Whole weight',  
'Whole weight.1', 'Whole weight.2', 'Shell weight', 'Rings']
```

Preprocessing: One-Hot Encode Sex and Align Columns

```
In [10]: # Separate target and features
y_train = train['Rings']
X_train = train.drop(columns=['Rings'])

# One-hot encode Sex using pandas.get_dummies
X_train_dum = pd.get_dummies(X_train, columns=['Sex'], drop_first=True)
X_test_dum = pd.get_dummies(test, columns=['Sex'], drop_first=True)

# Align columns in case some categories are missing in test
X_train_dum, X_test_dum = X_train_dum.align(X_test_dum, join='left', axis=1,
                                             fill_value=0)

print("Encoded train columns:", X_train_dum.columns.tolist())
```

Encoded train columns: ['id', 'Length', 'Diameter', 'Height', 'Whole weight', 'Whole weight.1', 'Whole weight.2', 'Shell weight', 'Sex_I', 'Sex_M']

Model 1: Regularized Regression (Ridge)

```
In [11]: # Standardize + RidgeCV
ridge_scaler = StandardScaler()
X_train_scaled = ridge_scaler.fit_transform(X_train_dum)
X_test_scaled = ridge_scaler.transform(X_test_dum)

ridge = RidgeCV(alphas=np.logspace(-3, 3, 50), cv=5)
ridge.fit(X_train_scaled, y_train)

ridge_pred = ridge.predict(X_test_scaled)
ridge_pred = np.maximum(ridge_pred, 0) # clip negatives

print("Ridge alpha selected:", ridge.alpha_)
print("First 10 Ridge predictions (clipped):")
print(ridge_pred[:10])

# Create Kaggle submission file for regularized model
sub_reg = pd.DataFrame({'id': test['id'], 'Rings': ridge_pred})
sub_reg.to_csv('abalone_regularized.csv', index=False)

print("\nSaved abalone_regularized.csv for Kaggle submission.")
```

Ridge alpha selected: 2.023589647725158
 First 10 Ridge predictions (clipped):
 [8.51308162 10.00665704 10.28233 10.93188897 7.75904457 9.63489889
 12.87044962 6.21137981 7.77861508 11.46983529]

Saved abalone_regularized.csv for Kaggle submission.

```
In [12]: # Diagnostics for Regularized Regression (Ridge) on training data

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

# Predict on training data
ridge_train_pred = ridge.predict(X_train_scaled)
```

```

y_true = y_train.values
y_pred = ridge_train_pred
resid_ridge = y_true - y_pred

# 1. Actual vs Predicted
plt.figure(figsize=(6, 4))
sns.scatterplot(x=y_true, y=y_pred, alpha=0.4)
lims = [min(y_true.min(), y_pred.min()), max(y_true.max(), y_pred.max())]
plt.plot(lims, lims, 'r--', linewidth=2)
plt.xlabel("Actual Rings")
plt.ylabel("Predicted Rings")
plt.title("Regularized Regression (Ridge): Actual vs Predicted")
plt.tight_layout()
plt.show()


# 2. Residuals vs Fitted
plt.figure(figsize=(6, 4))
sns.scatterplot(x=y_pred, y=resid_ridge, alpha=0.4)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Fitted Values (Predicted Rings)")
plt.ylabel("Residuals")
plt.title("Regularized Regression (Ridge): Residuals vs Fitted")
plt.tight_layout()
plt.show()

# 3. Residual Distribution with KDE
plt.figure(figsize=(6, 4))
sns.histplot(resid_ridge, kde=True)
plt.xlabel("Residual")
plt.title("Regularized Regression (Ridge): Residual Distribution with KDE")
plt.tight_layout()
plt.show()


# 4. Print basic error metrics on training data
rmse = np.sqrt(mean_squared_error(y_true, y_pred))
mae = mean_absolute_error(y_true, y_pred)
r2 = r2_score(y_true, y_pred)
print(f'Ridge Training RMSE: {rmse:.4f}')
print(f'Ridge Training MAE: {mae:.4f}')
print(f'Ridge Training R^2: {r2:.4f}')

```


Regularized Regression (Ridge): Actual vs Predicted



Regularized Regression (Ridge): Residuals vs Fitted



Regularized Regression (Ridge): Residual Distribution with KDE



```
Ridge Training RMSE: 1.9947  
Ridge Training MAE: 1.3739  
Ridge Training R^2: 0.6056
```

Model 2: Principal Components Regression (PCR)

```
In [13]: # Standardize for PCA  
pca_scaler = StandardScaler()  
X_train_scaled_pca = pca_scaler.fit_transform(X_train_dum)  
  
# Cross-validate number of components  
max_components = min(X_train_scaled_pca.shape[1], 20)  
cv = KFold(n_splits=5, shuffle=True, random_state=123)  
mse_per_comp = []  
  
for n_comp in range(1, max_components + 1):  
    pca = PCA(n_components=n_comp)  
    X_pca = pca.fit_transform(X_train_scaled_pca)  
    lm = LinearRegression()  
    scores = cross_val_score(lm, X_pca, y_train, cv=cv, scoring='neg_mean_squared_error')  
    mse = -scores.mean()  
    mse_per_comp.append((n_comp, mse))  
  
best_n_comp, best_mse = min(mse_per_comp, key=lambda x: x[1])  
print("Best number of components:", best_n_comp)  
print("Best CV MSE:", best_mse)  
  
# Fit final PCR model  
pca_final = PCA(n_components=best_n_comp)  
X_train_pca = pca_final.fit_transform(X_train_scaled_pca)  
lm_final = LinearRegression()
```

```

lm_final.fit(X_train_pca, y_train)

# Transform test data
X_test_scaled_pca = pca_scaler.transform(X_test_dum)
X_test_pca = pca_final.transform(X_test_scaled_pca)

pcr_pred = lm_final.predict(X_test_pca)
pcr_pred = np.maximum(pcr_pred, 0) # clip negatives

print("First 10 PCR predictions (clipped):")
print(pcr_pred[:10])

# Create Kaggle submission file for PCR model
sub_pcr = pd.DataFrame({'id': test['id'], 'Rings': pcr_pred})
sub_pcr.to_csv('abalone_pcr.csv', index=False)

print("\nSaved abalone_pcr.csv for Kaggle submission.")

```

Best number of components: 10
 Best CV MSE: 3.981227439595023
 First 10 PCR predictions (clipped):
 [8.51178779 10.00607823 10.28222843 10.93256938 7.75913987 9.63463772
 12.87092589 6.21136134 7.77853918 11.46874172]

Saved abalone_pcr.csv for Kaggle submission.

In [14]: # Diagnostics for Principal Components Regression (PCR) on training data

```

import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

# Reconstruct the training set in the selected PCA space
# (Assumes pca_scaler, pca_final, and lm_final were fit above)
X_train_scaled_pca = pca_scaler.transform(X_train_dum)
X_train_pca_final = pca_final.transform(X_train_scaled_pca)

pcr_train_pred = lm_final.predict(X_train_pca_final)
y_true = y_train.values
y_pred = pcr_train_pred
resid_pcr = y_true - y_pred

# 1. Actual vs Predicted
plt.figure(figsize=(6, 4))
sns.scatterplot(x=y_true, y=y_pred, alpha=0.4)
lims = [min(y_true.min(), y_pred.min()), max(y_true.max(), y_pred.max())]
plt.plot(lims, lims, 'r--', linewidth=2)
plt.xlabel("Actual Rings")
plt.ylabel("Predicted Rings")
plt.title("PCR: Actual vs Predicted")
plt.tight_layout()
plt.show()

# 2. Residuals vs Fitted
plt.figure(figsize=(6, 4))

```


```

sns.scatterplot(x=y_pred, y=resid_pcr, alpha=0.4)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel("Fitted Values (Predicted Rings)")
plt.ylabel("Residuals")
plt.title("PCR: Residuals vs Fitted")
plt.tight_layout()
plt.show()


# 3. Residual Distribution with KDE
plt.figure(figsize=(6, 4))
sns.histplot(resid_pcr, kde=True)
plt.xlabel("Residual")
plt.title("PCR: Residual Distribution with KDE")
plt.tight_layout()
plt.show()

# 4. Print basic error metrics on training data
rmse = np.sqrt(mean_squared_error(y_true, y_pred))
mae = mean_absolute_error(y_true, y_pred)
r2 = r2_score(y_true, y_pred)
print(f"PCR Training RMSE: {rmse:.4f}")
print(f"PCR Training MAE: {mae:.4f}")
print(f"PCR Training R^2: {r2:.4f}")


```



PCR: Residuals vs Fitted



PCR: Residual Distribution with KDE



PCR Training RMSE: 1.9947

PCR Training MAE: 1.3739

PCR Training R²: 0.6056